

Sequence to Sequence Models

RECURRENT NEURAL NETWORKS FOR LANGUAGE MODELING IN PYTHON



David Cecchini
Instructor

Sequence to sequence

Possible architectures:

- Many inputs with one output
 - Sentiment analysis
 - Classification
- Many inputs to many outputs
 - Text generation
 - Neural Machine Translation (NMT)

Text generation: example

Text generation: example

```
# Pre-trained model  
model.generate_sheldon_phrase()
```

```
'knock knock. penny. do you have an epost is part in your expert,  
too bealie to play the tariment with last night.'
```

Text generation: modeling

How to build text generation models:

- Decide if a token will be characters or words
 - Words demands very large datasets (hundred of millions sentences)
 - Chars can be trained faster, but can generate typos
- Prepare the data
 - Build training sample with (past tokens, next token) examples
- Design the model architecture
 - Embedding layer, number of layers, etc.
- Train and experiment

NMT: example

Neural Machine Translation: example

```
# Pre-trained model  
model.translate("Vamos jogar futebol?")
```

```
'Let's go play soccer?'
```

NMT: modeling

How to build NMT models:

- Get a sample of translated sentences
 - For example, the **Anki project**
- Prepare the data
 - Tokenize input language sentences
 - Tokenize output language sentences
- Design the model architecture
 - Encoder and decoder
- Train and experiment

Chapter outline

In this chapter:

- Text Generation
 - Use pre-trained model to generate a sentence
 - Learn to prepare the data and build the model
- Neural Machine Translation (NMT)
 - All-in-one NMT model

Let's practice!

RECURRENT NEURAL NETWORKS FOR LANGUAGE MODELING IN PYTHON

The Text Generating Function

RECURRENT NEURAL NETWORKS FOR LANGUAGE MODELING IN PYTHON



David Cecchini
Data Scientist

Generating sentences

- Sentence is determined by punctuation. For example, `.` (period), `!` (exclamation) or `?` (question).
 - The punctuation marks need to be in the vocabulary.
- There is a sentence token, e.g. `<SENT>` and `</SENT>`, that determines when a sentence begins and ends.
 - Need to pre-process the data to insert the labels.

Generating sentences

```
sentence = ''  
  
# Loop until end of sentence  
while next_char != '.':  
    # Predict next char: Get pred array in position 0  
    pred = model.predict(X)[0]  
    char_index = np.argmax(pred)  
    next_char = index_to_char(char_index)  
    # Concatenate to sentence  
    sentence = sentence + next_char
```

Probability scaling

Scale the probability distribution.

- **Temperature:** name from physics
 - *Small values:* makes prediction more confident
 - *Value equal to one:* no scaling
 - *higher values:* makes prediction more creative
 - Hyper-parameter: Try different values to fit the predictions to your need

Probability scaling

```
def scale_softmax(softmax_pred, temperature=1.0):  
    # Take the logarithm  
    scaled_pred = np.log(softmax_pred) / temperature  
  
    # Re-apply the exponential  
    scaled_pred = np.exp(scaled_pred)  
  
    # Build probability distribution  
    scaled_pred = scaled_pred / np.sum(scaled_pred)  
  
    # Simulate multinomial  
    scaled_pred = np.random.multinomial(1, scaled_pred, 1)  
  
    # Return simulated class  
    return np.argmax(scaled_pred)
```

Let's practice!

RECURRENT NEURAL NETWORKS FOR LANGUAGE MODELING IN PYTHON

Text Generation Models

RECURRENT NEURAL NETWORKS FOR LANGUAGE MODELING IN PYTHON



David Cecchini
Data Scientist

Similar to a classification model

The Text Generation Model:

- Uses the vocabulary as classes
- The last layer applies a softmax with vocabulary size units
- Uses `categorical_crossentropy` as loss function

Example model using keras

```
model = Sequential()

model.add(LSTM(units, input_shape=(chars_window, n_vocab),
              dropout=0.15, recurrent_dropout=0.15, return_sequences=True))

model.add(LSTM(units, dropout=dropout, recurrent_dropout=0.15,
              return_sequences=False))

model.add(Dense(n_vocab, activation='softmax'))

model.compile(loss='categorical_crossentropy', optimizer='adam')
```

But not really classification model

Difference to classification:

- Computes loss, but not performance metrics (accuracy)
 - Humans see results and evaluate performance.
 - If not good, train more epochs or add complexity to the model (add more memory cells, add layers, etc.).
- Used with generation rules according to task
 - Generate next char
 - Generate one word
 - Generate one sentence
 - Generate one paragraph

Other applications

- Name creation
 - Baby names
 - New star names, etc.
- Generate marked text
 - LaTeX
 - Markdown
 - XML, etc.
 - Programming code
- News articles
- Chatbots

Data prep

I am not insane,
my mother had
me tested

step = 1



Sentences	Next char
I	\b
I\b	a
I a	m
I am	\b



chars_window = 10
vocabulary = 5

X	Y
[0 0 0 0 1 0 0 0 0 0]	[0 1 0 0 0]
[0 1 0 0 1 0 0 0 0 0]	[1 0 0 0 0]
[0 1 1 0 1 0 0 0 0 0]	[0 0 0 0 1]
[0 1 1 0 1 0 0 1 0 0]	[0 1 0 0 0]

Let's practice!

RECURRENT NEURAL NETWORKS FOR LANGUAGE MODELING IN PYTHON

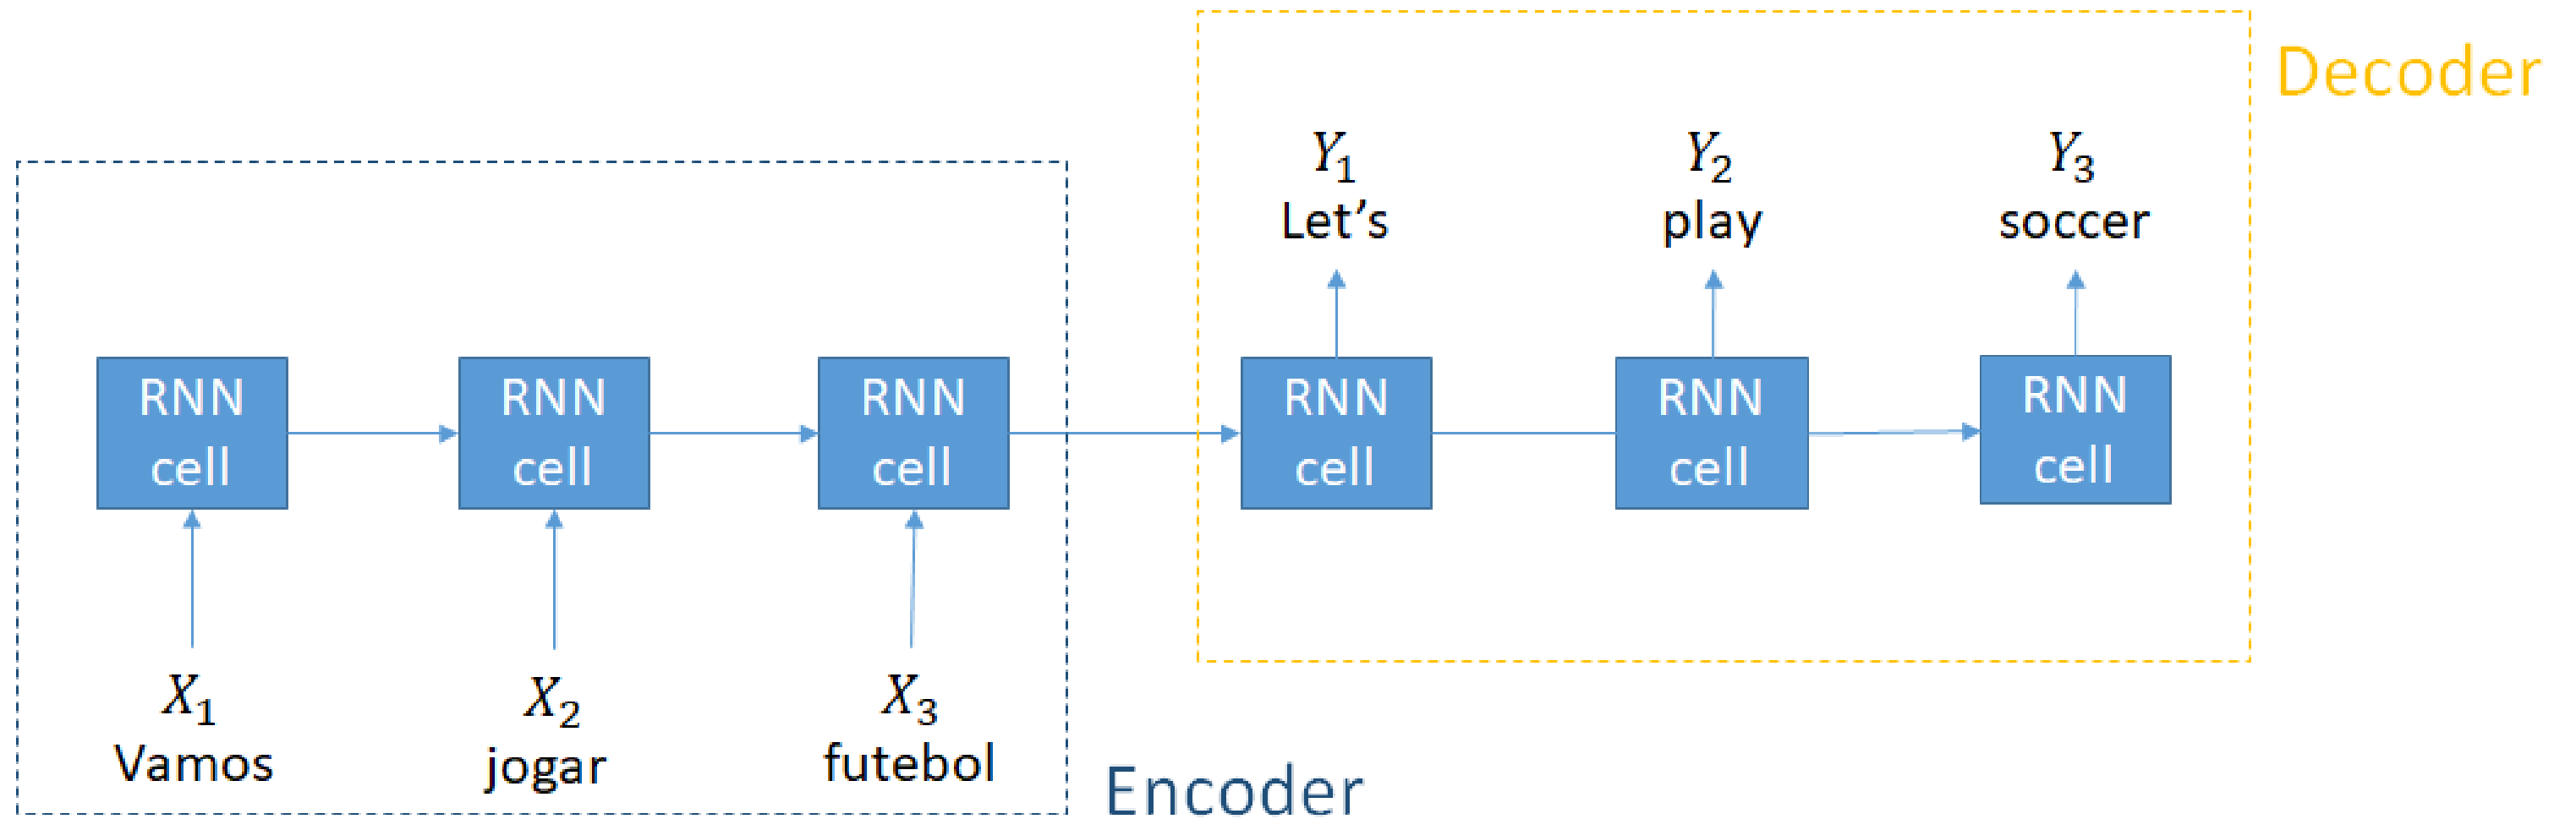
Neural Machine Translation

RECURRENT NEURAL NETWORKS FOR LANGUAGE MODELING IN PYTHON



David Cecchini
Data Scientist

Encoder and decoders



Encoder example

```
# Instantiate the model
model = Sequential()

# Embedding layer for input language
model.add(Embedding(input_language_size, input_wordvec_dim,
                    input_length=input_language_len, mask_zero=True))

# Add LSTM layer
model.add(LSTM(128))

# Repeat the last vector
model.add(RepeatVector(output_language_len))
```



Decoder example

```
# Right after the encoder
model.add(LSTM(128, return_sequences=True))

# Add Time Distributed
model.add(TimeDistributed(Dense(eng_vocab_size, activation='softmax')))
```


Data prep

Vamos jogar futebol esse domingo



[2, 5, 12, 10, 15]

Let's play soccer this Sunday



[11, 3, 15, 10, 7]

Vocab size
= 15



0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0
0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1
0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0
0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0

Data preparation for the input language

```
# Import modules
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
```

```
# Use the Tokenizer class
tokenizer = Tokenizer()
tokenizer.fit_on_texts(input_texts_list)

# Text to sequence of numerical indexes
X = tokenizer.texts_to_sequences(input_texts_list)

# Pad sequences
X = pad_sequences(X, maxlen=length, padding='post')
```

Tokenize the output language

```
# Use the Tokenizer class
tokenizer = Tokenizer()
tokenizer.fit_on_texts(output_texts_list)

# Text to sequence of numerical indexes
Y = tokenizer.texts_to_sequences(output_texts_list)

# Pad sequences
Y = pad_sequences(Y, maxlen=length, padding='post')
```

One-hot encode the output language

```
# Instantiate a temporary variable
ylist = list()

# Loop over the sequence of numerical indexes
for sequence in Y:

    # One-hot encode each index on current sentence
    encoded = to_categorical(sequence, num_classes=vocab_size)

    # Append one-hot encoded values to the list
    ylist.append(encoded)

# Transform to np.array and reshape
Y = np.array(ylist).reshape(Y.shape[0], Y.shape[1], vocab_size)
```

Note on training and evaluating

Training the model:

```
model.fit(X, Y, epochs=N)
```

Evaluating:

- Use BLEU
 - `nltk.translate.bleu_score`

Let's practice!

RECURRENT NEURAL NETWORKS FOR LANGUAGE MODELING IN PYTHON

Congratulations!

RECURRENT NEURAL NETWORKS FOR LANGUAGE MODELING IN PYTHON



David Cecchini
Data Scientist

Wrap-up

- Introduction to language tasks:
 - Sentiment classification
 - Multi-class classification
 - Text Generation
 - Neural Machine Translation
- Sequence to sequence models
- Implementation in Keras

RNN pitfalls and different cell types

- Vanishing and exploding gradient problems
- GRU and LSTM cells
- Word vectors and the Embedding layer
- Better sentiment analysis

Multi-class classification

- Data preparation
- Transfer learning
- Keras models
- Model performance

Text generation and NMT

- Text Generation
 - Chars as token
 - Data preparation
 - Generate sentences mimicking Sheldon
- Neural Machine Translation
 - Words as tokens
 - Data preparation: encoders and decoders
 - Translate Portuguese to English

Congratulations!!!

RECURRENT NEURAL NETWORKS FOR LANGUAGE MODELING IN PYTHON